Coursework 1 - Modelling

Jonas Osborn jo14944

Tristan Saunders ts16802

Corin Varney cv14985

1 The prior

1.1 Theory

Question 1 *1.* As the instances of y are noisy observations of the underlying process and we do not know anything about this uncertainty we can assume it is the sum of independent and identically distributed errors. The Central Limit Theorem states the distribution of the sum of a large enough number of independent, identically distributed variables will be approximately normally distributed. From this we can say that our model of y has the following form:

$$y = f(x) + \epsilon$$
$$\text{where: } \epsilon \sim \mathcal{N}(0, \boldsymbol{I})$$

And so from this we have the likelihood of each y_i as a Gaussian distribution with mean $f(x_i)$.

2. Choosing a spherical covariance matrix for the likelihood means that we are assuming the different dimensions of y_i to be independent and identically distributed. As they are independent they do not covary with each other and so the covariance matrix is diagonal. As they are identically distributed they all have the same variance. Therefore the covariance matrix is spherical.

Question 2

$$p(\boldsymbol{Y}|f, \boldsymbol{X}) = p(\boldsymbol{y}_N|\boldsymbol{y}_{N-1}, \ldots, \boldsymbol{y}_1, f, \boldsymbol{X})p(\boldsymbol{y}_{N-1}|\boldsymbol{y}_{N-2}, \ldots, \boldsymbol{y}_1, f, \boldsymbol{X}) \ldots p(\boldsymbol{y}_1|f, \boldsymbol{X})$$

1.1.1 Linear regression

Question 3

$$p(\boldsymbol{Y}|\boldsymbol{X}, \boldsymbol{W}) = \prod_i^N p(\boldsymbol{y}_i|\boldsymbol{x}_i, \boldsymbol{W})$$
$$p(\boldsymbol{Y}|\boldsymbol{X}, \boldsymbol{W}) = \prod_i^N \mathcal{N}(\boldsymbol{W}\boldsymbol{x}_i, \sigma^2\boldsymbol{I})$$

Question 4 A conjugate prior is one that is conjugate to the posterior, meaning they are in the same family of distributions. A conjugate prior is useful as it gives a closed-form solution for the posterior. If we didn't choose a conjugate prior then numerical integration may be necessary to calculate the posterior which may mean the solution is potentially intractable. The conjugate prior for a Gaussian posterior is a Gaussian.

Question 5 todo

Question 6 todo

1.1.2 Non-parametric regression

Question 7 todo

Question 8 todo

Question 9 todo

Question 10 todo

Question 11 todo

1.2 Practical

1.2.1 Linear regression

Question 12 todo

1.2.2 Non-parametric regression

Question 13 todo

Question 14 todo

2 The posterior

2.1 Theory

2.1.1 Learning

2.1.2 Practical optimisation

2.1.3 Non-parametric

2.2 Practical

2.2.1 Linear representation learning